

Tech Community Sharing

# Azure OpenAI 集成 云原生应用



# Overview

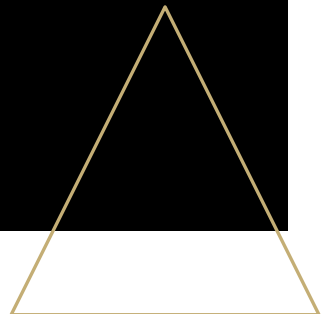
Architecture

Semantic Kernel

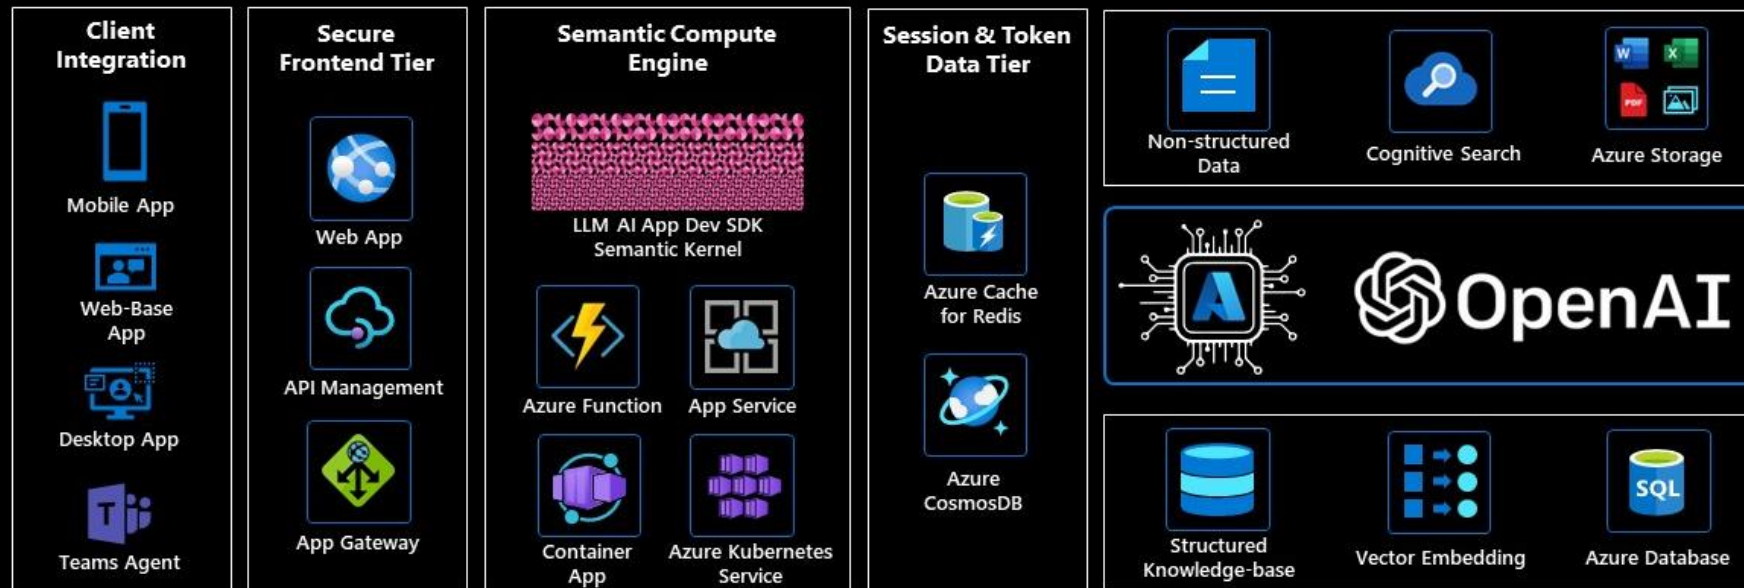
Q & A

Details

Demo



# Enterprise GPT-size Intelligent Application Architecture



A prompt engineering approach with LLM AI Dev Framework



Azure Cognitive Search can quickly index unstructured data such as PDF and WORD files, allowing existing data to be used immediately.



Utilizing Azure Database with vector storage and processing capabilities and combining them with AOAI's embedding vector generation model, the enterprise's existing structured knowledge base can be integrated easily

## Dual-engine



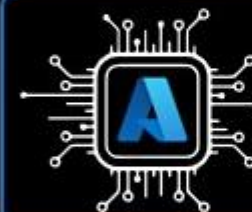
Non-structured  
Data



Cognitive Search



Azure Storage



 OpenAI



Structured  
Knowledge-base

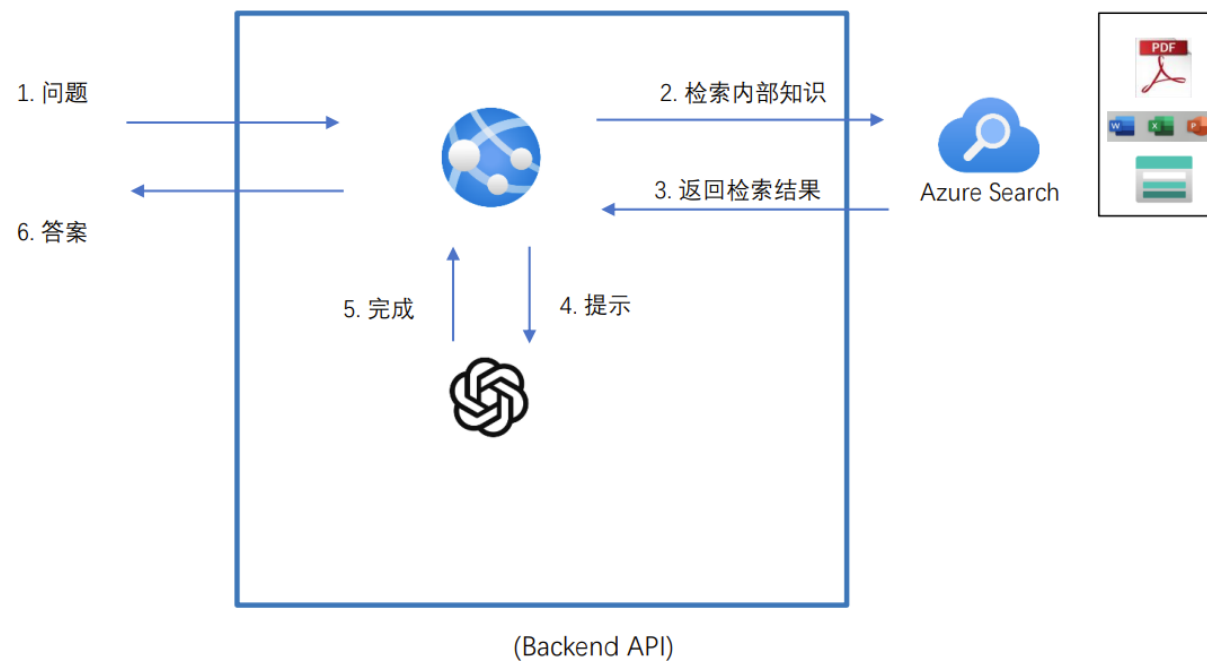


Vector Embedding



Azure Database

# GPT + enterprise knowledge base/data



# Azure Cognitive Search

ingest

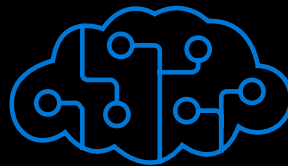
abundant

Browse

101010  
010101  
101010

Data in any  
format, any  
Azure store

Cognitive skills

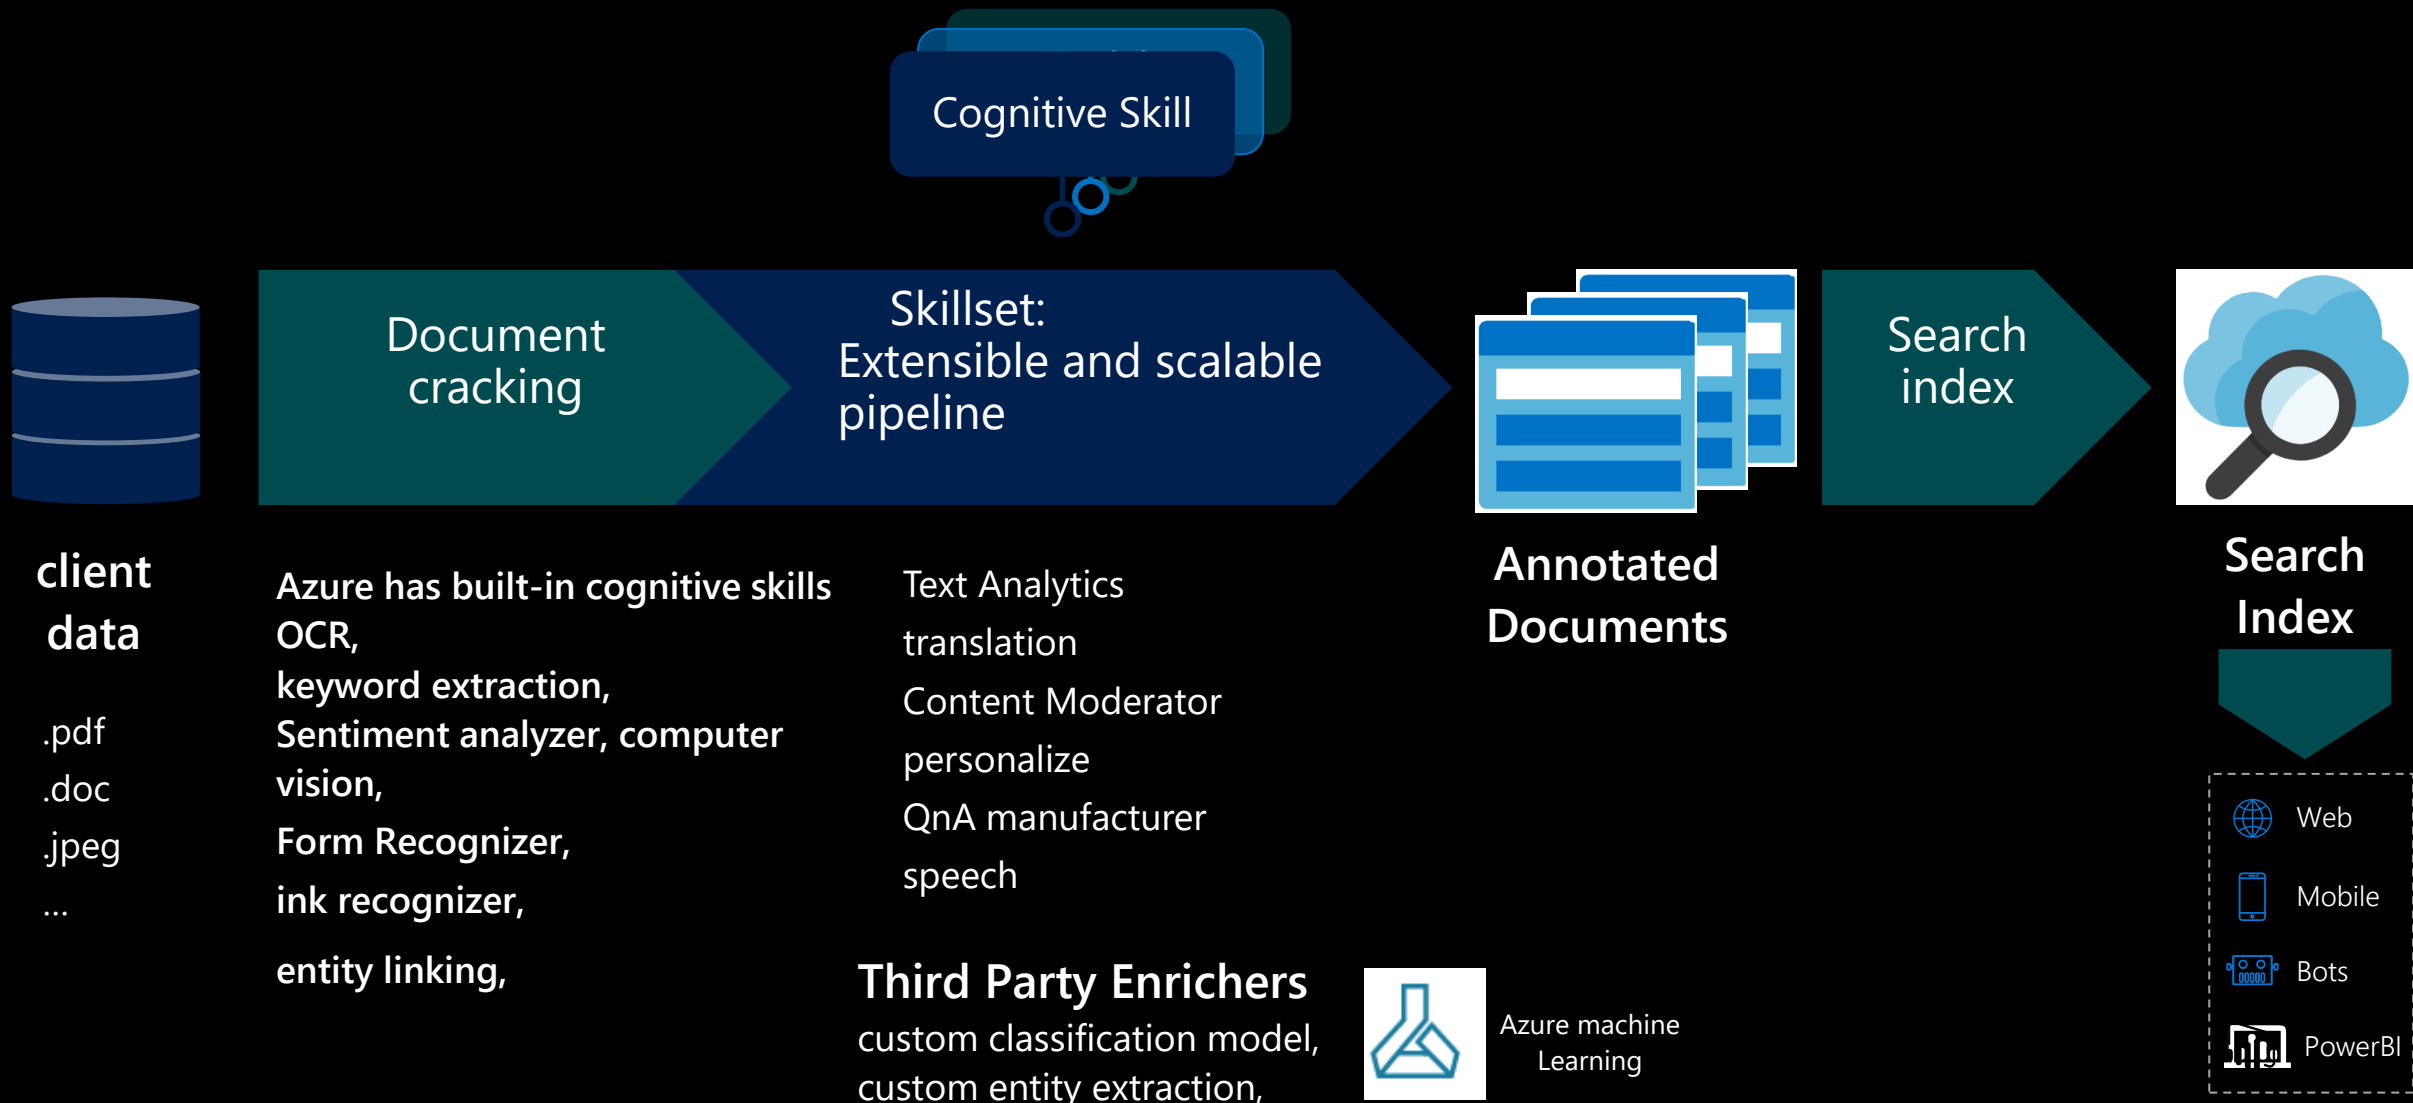


Search



Annotations

# Cognitive Search Architecture



### Session & Token Data Tier



Azure Cache  
for Redis

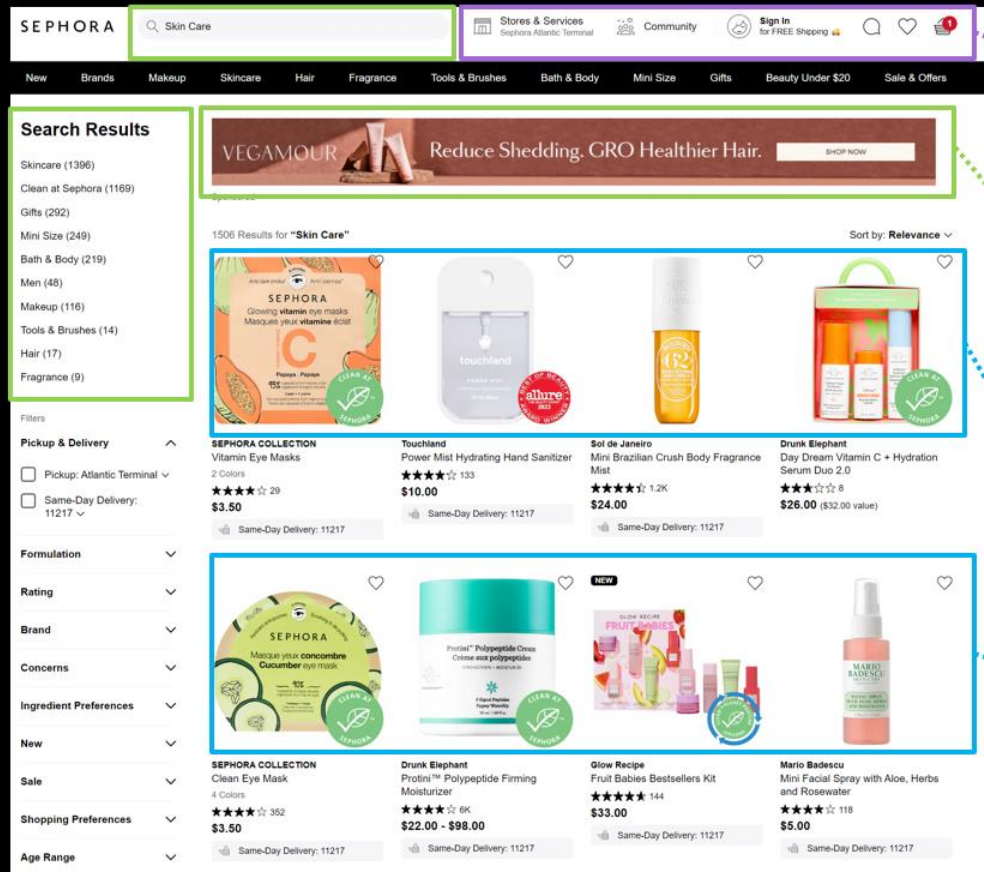


Azure  
CosmosDB

- Based on Redis and CosmosDB.
- Add contextual caching, session persistence, prompt persistence, and other capabilities to the application.
- Leave room for future model or engine optimization based on prompts.



# How is Azure Cache for Redis Leveraged?



## Session Store

Easily resume a user's journey on any mobile or web application by quickly sorting the data of their last session regardless of where they are across the globe



## Message Broker Store

Azure Cache for Redis can communicate across all Microservices. In the example, it communicated with its ad server to give a recommended ad for the shopper



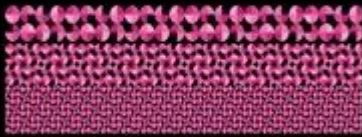
## Distributed Cache

All images need to be consistently accessed and are static therefore, minimizes load times by keeping all assets into a cache no matter where the user is across the globe



- The API encapsulation, load balancing, and gateway at the front end further enhance application security and reliability.
- Allow the intelligent entity to integrate with various front-end apps in a more secure and stable manner.

### Semantic Compute Engine



LLM AI App Dev SDK  
Semantic Kernel



Azure Function



App Service



Container  
App



Azure Kubernetes  
Service

- Utilize a flexible semantic compute engine as prompt engine
- Support various deployment forms such as PaaS, Serverless, and containers
- Enable enterprises to optimize the prompt engine with modern LLM AI application development framework
- Elastically scale computing resources according to real-time business needs.

# Function as a Service

Use serverless code to handle events.

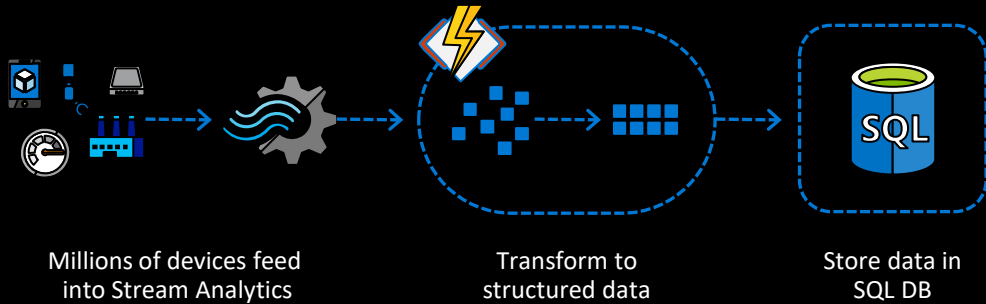
- Writing cloud apps made easy
- Extend functionality based on customer needs
- Develop functions in languages such as C#, Node.js, F#, Python, Java, and more
- Easily schedule event-driven tasks across services
- Exposes the function as an HTTP API endpoint



# Azure Functions Common application scenarios

Any scenario that requires a response to an event

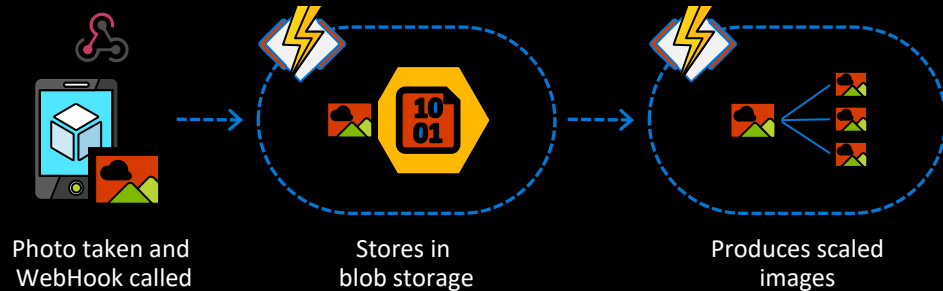
## Real-time stream processing



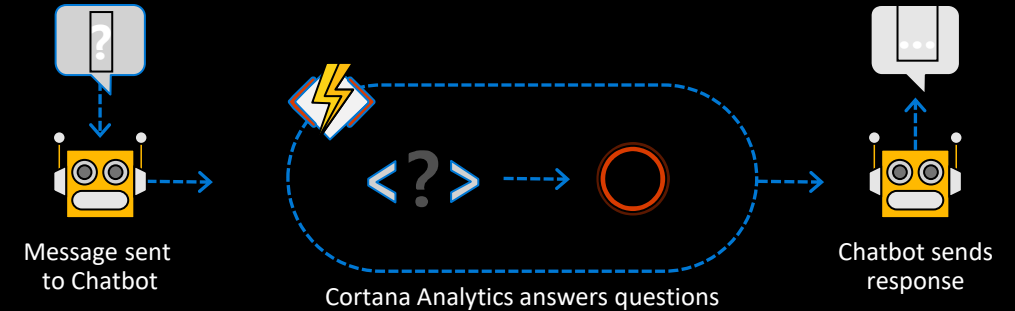
## Timer-based processing



## Backend (Mobile/IoT/Web)



## Real-time bot messaging

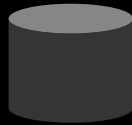


# Semantic Kernel - LLM AI Dev Framework

## THE CORE

Everything you need to manage complex prompts, chains, long-running tasks, and **planning**.

SKILLS



PLANNER



RUNTIME

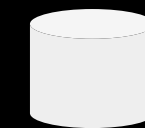


MEMORIES

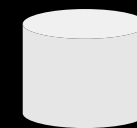
CONNECTORS

## THE BFD

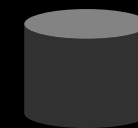
But wait, there's more! All the longform **memories** readily available in the MS Graph. Plus, the 900 Power **connectors**.



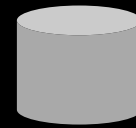
SUBSTRATE



LINKEDIN



GITHUB



POWER  
PLATFORM

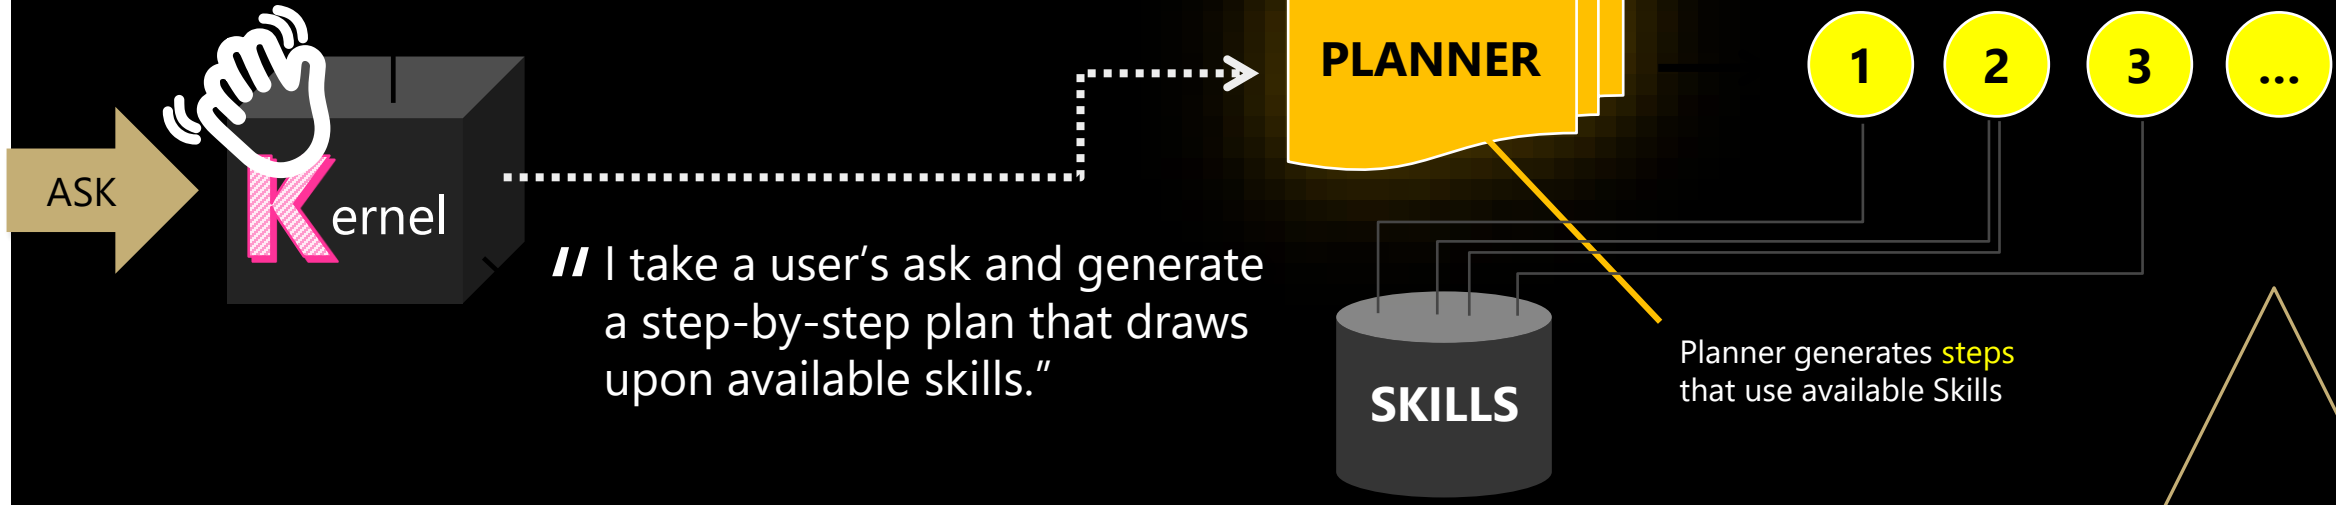
# Meet the lightweight Kernel of Semantic Kernel

```
using Microsoft.SemanticKernel;
```

```
var myKernel = Kernel.Builder.Build();
```

// I've been designed to reduce hallucinations, orchestrate complicated LLM AI prompts combined with native code, use multiple AI models, and ... I have a special skill to **PLAN**."

// I take a user's ask and generate a step-by-step plan that draws upon available skills."



# SK R1 Tour

Available NOW

It all starts with  
a user's AI ask ...

ASK



Steps ready  
from planner

1 2 3 ...

UNDER CONSTRUCTION

SUBSTRATE

LINKEDIN

GITHUB

Gather  
Skills

Gather  
Contexts

Gather  
Connectors

UNDER CONSTRUCTION

POWER  
PLATFORM

V1 READY  
SKILLS

RUNNING STEPS PIPELINE

GET

Result  
is ready

<

< 1

< 2

< 3

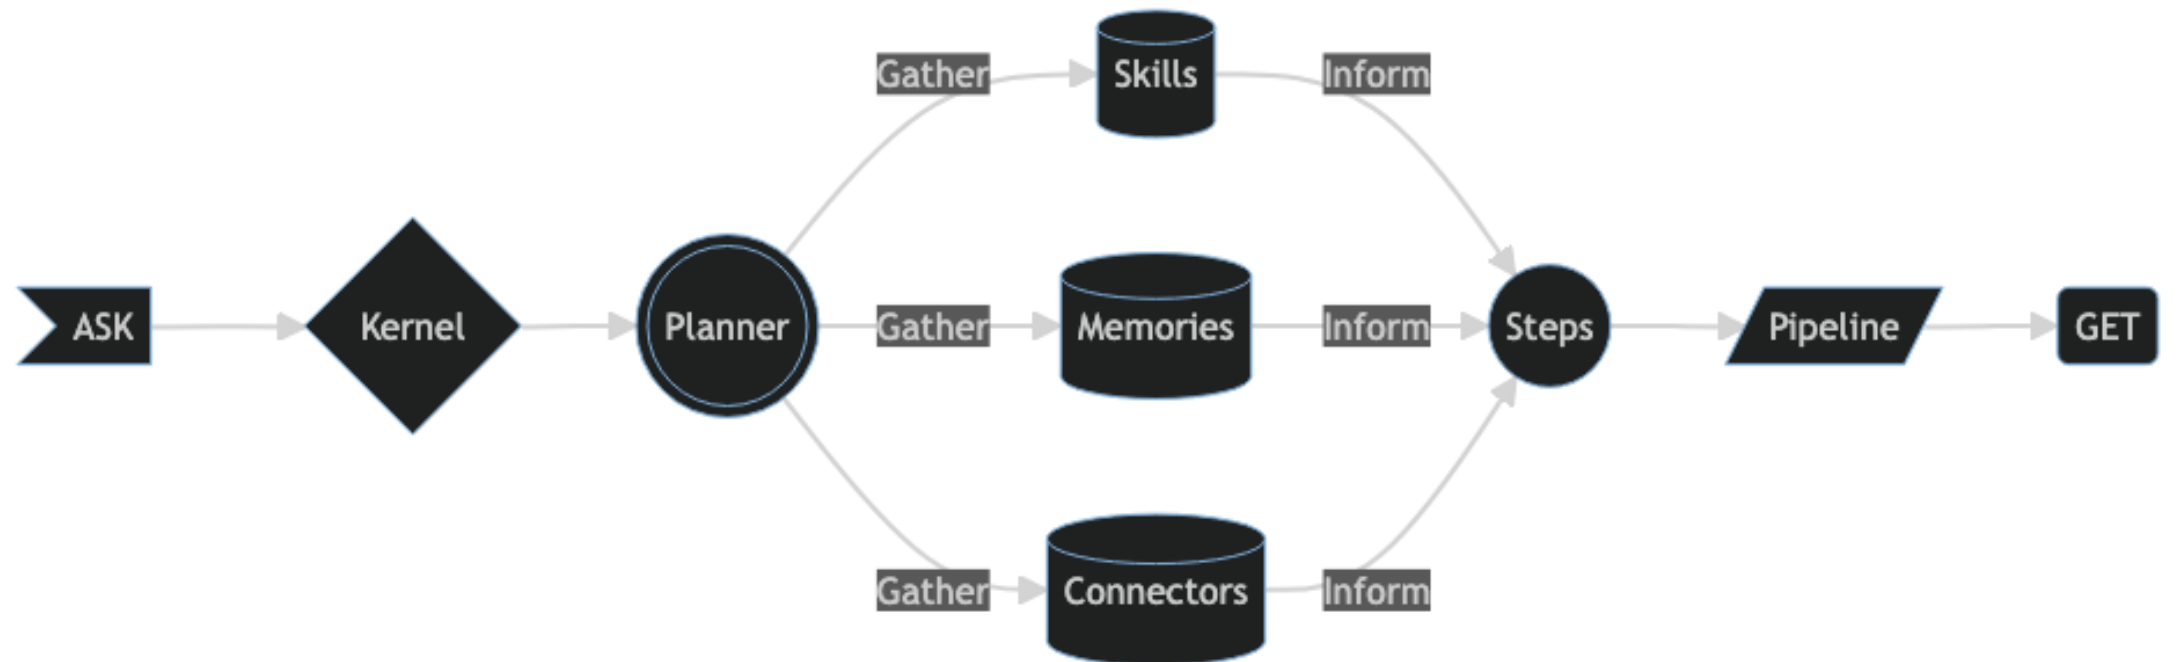
< ...

Execute  
Steps

... resulting in  
new productivity

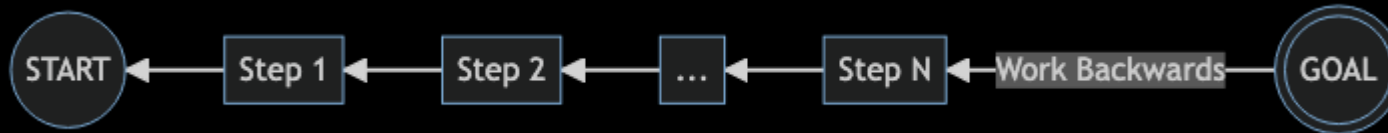


# A more "goal-oriented" approach to problem solving



# Planner

- The planner will operate within the skills it has available. In the event that a desired skill does not exist, the planner can suggest you create the skill. Or, depending upon the level of complexity the kernel can help you write the missing skill.



# Skill

- an LLM AI prompt — also called a "semantic" function
- native computer code -- also called a "native" function

SkillName (directory name)

└── Function1Name (directory name) skprompt.txt  
config.json

└── Function2Name (directory name)

MyAppSource

└── MySkillsDirectory

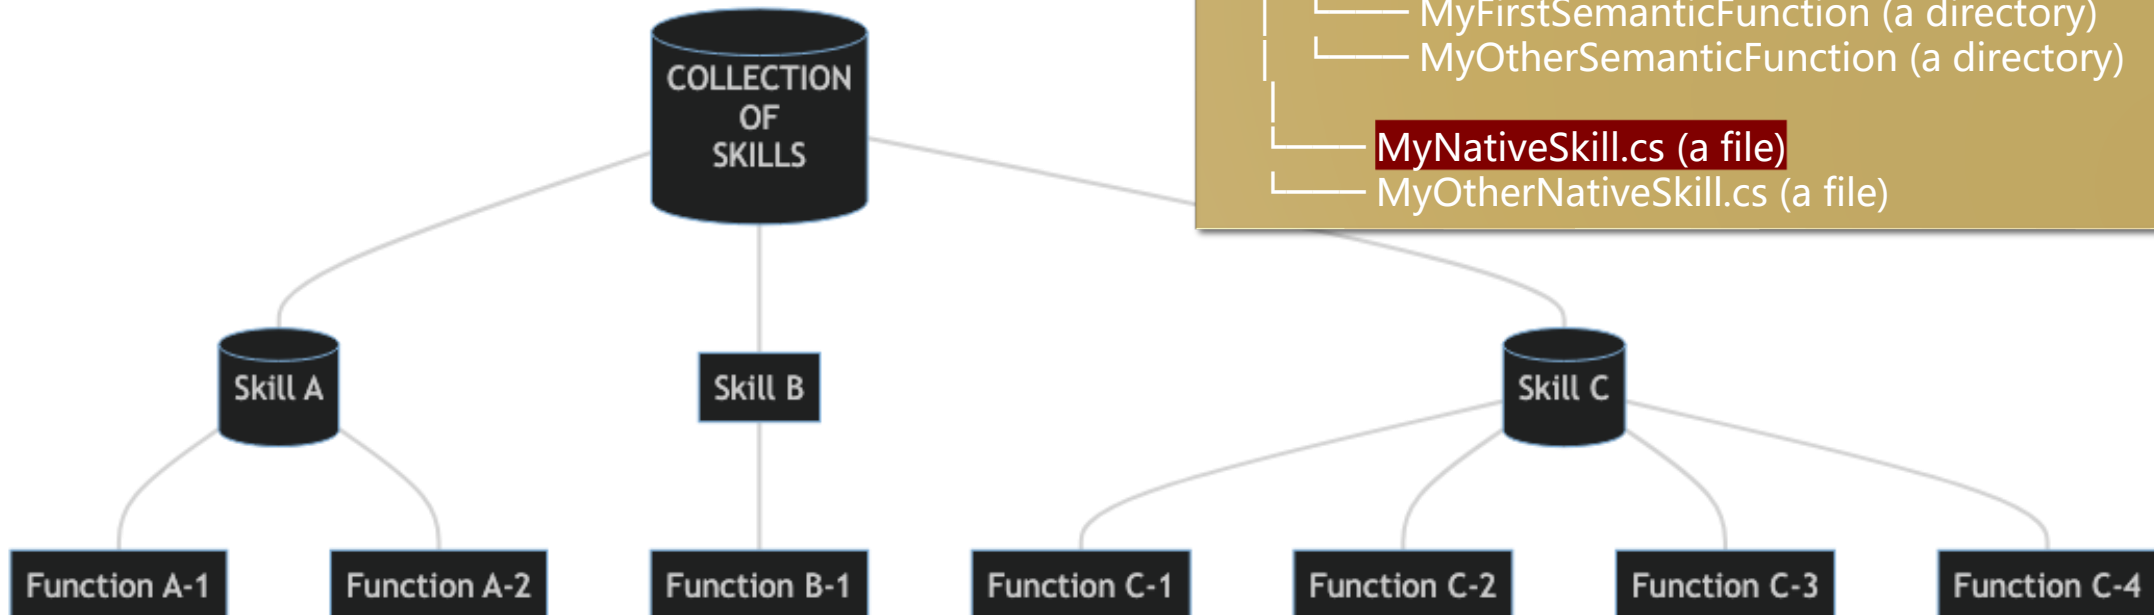
└── MySemanticSkill (a directory)

└── MyFirstSemanticFunction (a directory)

└── MyOtherSemanticFunction (a directory)

└── **MyNativeSkill.cs (a file)**

└── MyOtherNativeSkill.cs (a file)



# Memory

**Conventional key-value pairs:** Just like you would set an environment variable in your shell, the same can be done when using SK. The lookup is "conventional" because it's a one-to-one match between a key and your query.

**Conventional local-storage:** When you save information to a file, it can be retrieved with its filename. When you have a lot of information to store in a key-value pair, you're best off keeping it on disk.

**Semantic memory search:** You can also represent text information as a long vector of numbers, known as "embeddings." This lets you execute a "semantic" search that compares meaning-to-meaning with your query.

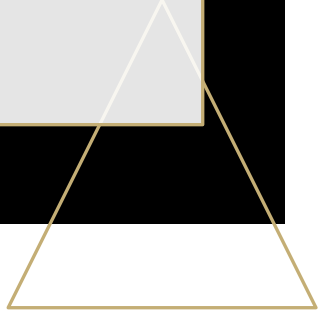


# Connector

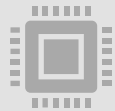
## MS Graph Connector Kit

- Add an event to your calendar
- Send an email for you
- Add a file to your OneDrive
- Create a share link to a file in your OneDrive
- Query your organization hierarchy
- Manage your MS To Do list

## Out-of-box

- Issue a Bing search query
  - Read OpenXML streams (e.g. Word docs)
  - Use SQLite as a lightweight database
- 

# SK makes app developers' work lives easier



**Fast Integration:** SK is designed to be embedded in any kind of application, making it easy for developers to add LLM AI functionality to test inside their apps.



**Power Prompting:** Plain prompts that are fed as API calls can only get you so far. SK provides the abstractions and machinery to unlock your OpenAI or Azure OpenAI API key.



**Novel-But-Familiar:** For 100% determininism, native code is always available as a first-class partner on your prompt engineering quests. You get the best of both worlds.



# How to use Semantic Kernel R1 in just 1 minute



## 1 MINUTE

Install the nuget package and go

```
#r "nuget: Microsoft.SemanticKernel, *-*"
```



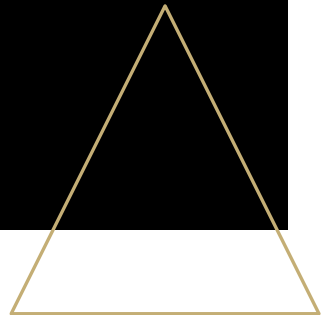
## HAVE MORE TIME?

Go deeper with  
the GitHub repo  
[aka.ms/skrepo](https://aka.ms/skrepo)



## MORE MINUTES

Learn more about  
its history  
[aka.ms/sk](https://aka.ms/sk)





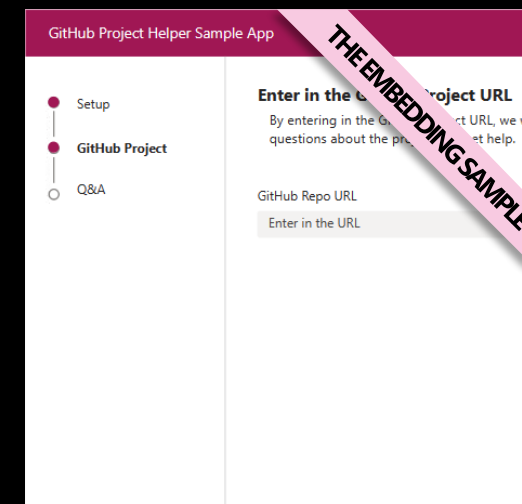
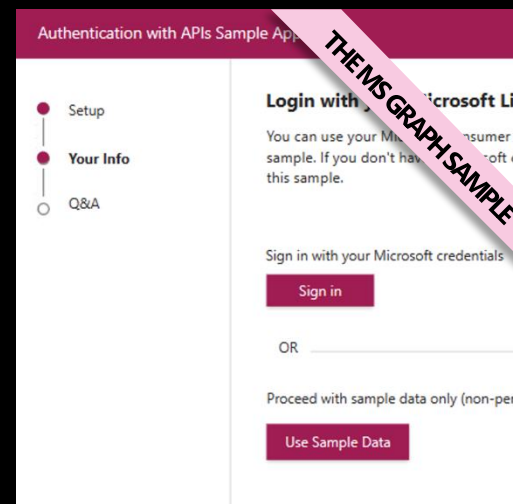
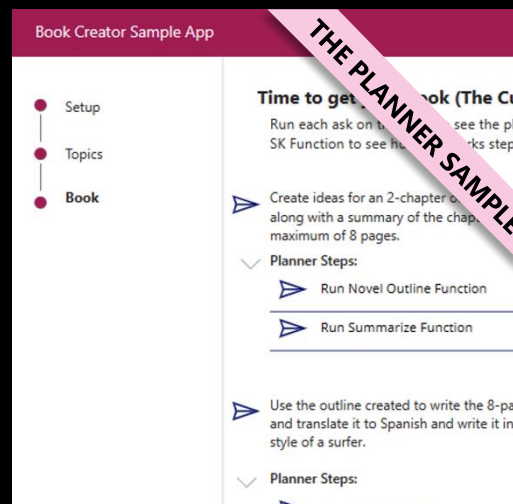
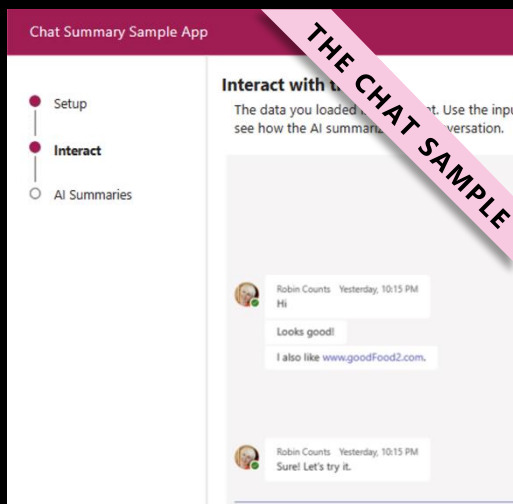
# The growing set of samples are designed for most use cases.

Easily add chat-themed AI skills into your app.

Chain prompts serially or hierarchically via Planner.

Query MS Graph information with your favorite prompts.

Chat with your repo of choice using embeddings.







# Demo

[SK Demos \(sharepoint.com\)](https://sharepoint.com)

[microsoft/semantic-kernel: Integrate cutting-edge LLM technology quickly and easily into your apps \(github.com\)](https://microsoft.com/semantic-kernel)

# Semantic Kernel Service API

semantic-kernel/samples/dotnet/KernelHttpServer

>>func start -csharp

[2023-04-02T13:18:19.291Z] Found C:\Users\huolu\Desktop\semantic-kernel\samples\dotnet\KernelHttpServer\KernelHttpServer.csproj. Using for user secrets file configuration.

Functions:

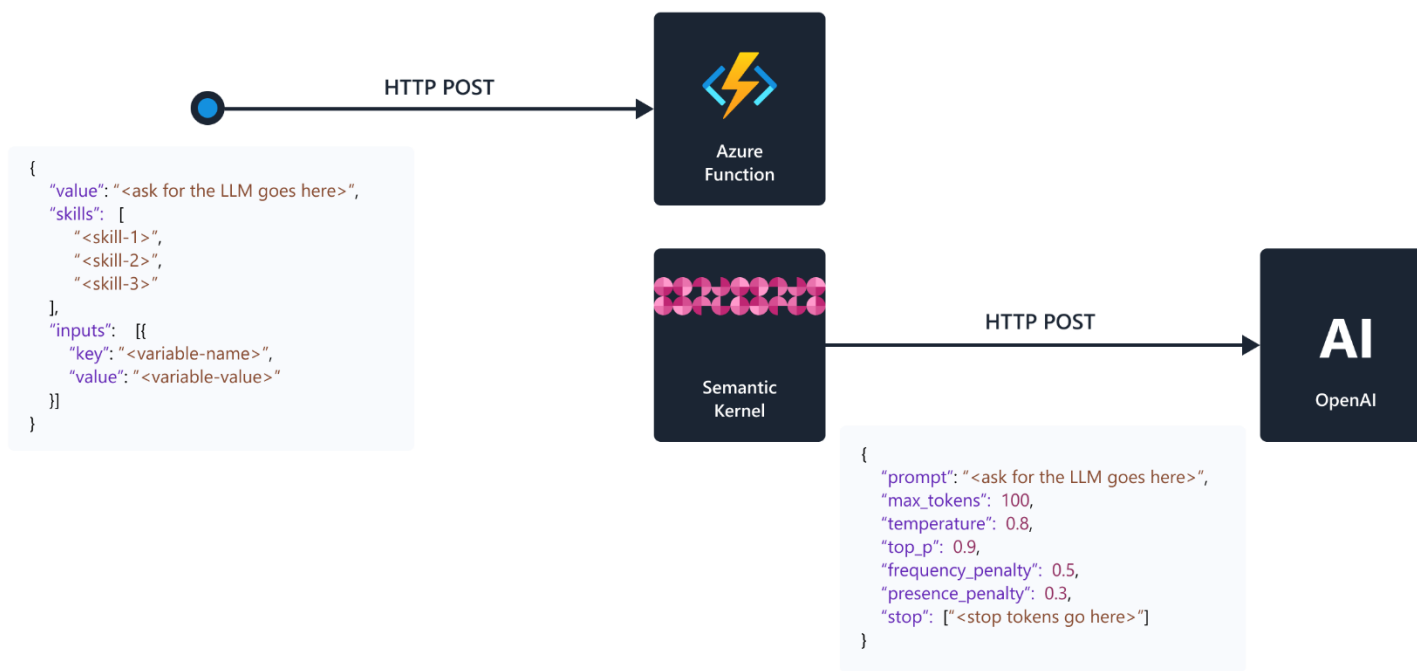
ExecutePlan: [POST] http://localhost:7071/api/planner/execute/{maxSteps?}

InvokeFunction: [POST] http://localhost:7071/api/skills/{skillName}/invoke/{functionName}

Ping: [GET] http://localhost:7071/api/ping

- Written in C# against Azure Function Runtime v4
- Expose some Semantic Kernel APIs that you can call via HTTP POST requests

Azure Functions Core Tools



# GitHub Repo Q&A Bot

```
semantic-kernel/samples/apps/github-qna-webapp-react
```

```
>>yarn install
```

```
>>yarn start
```

```
Compiled successfully!
```

```
You can now view starter-github-qna-webapp-react in the browser.
```

```
http://localhost:3000
```

```
Note that the development build is not optimized.  
To create a production build, use yarn build.
```

```
webpack compiled successfully  
No issues found.
```

The screenshot displays the GitHub Repo Q&A Bot web application. The browser address bar shows `localhost:3000`. The page title is "GitHub Repo Q&A Bot".

**Left Sidebar:**

- Setup
- Completion Model
- Embedding Model
- GitHub Repository
- Q&A**

**Main Content Area:**

### Ask questions about the repository

Now that the repository is in your local memory and embeddings are created for it, you can ask questions about it.

**Chat History:**

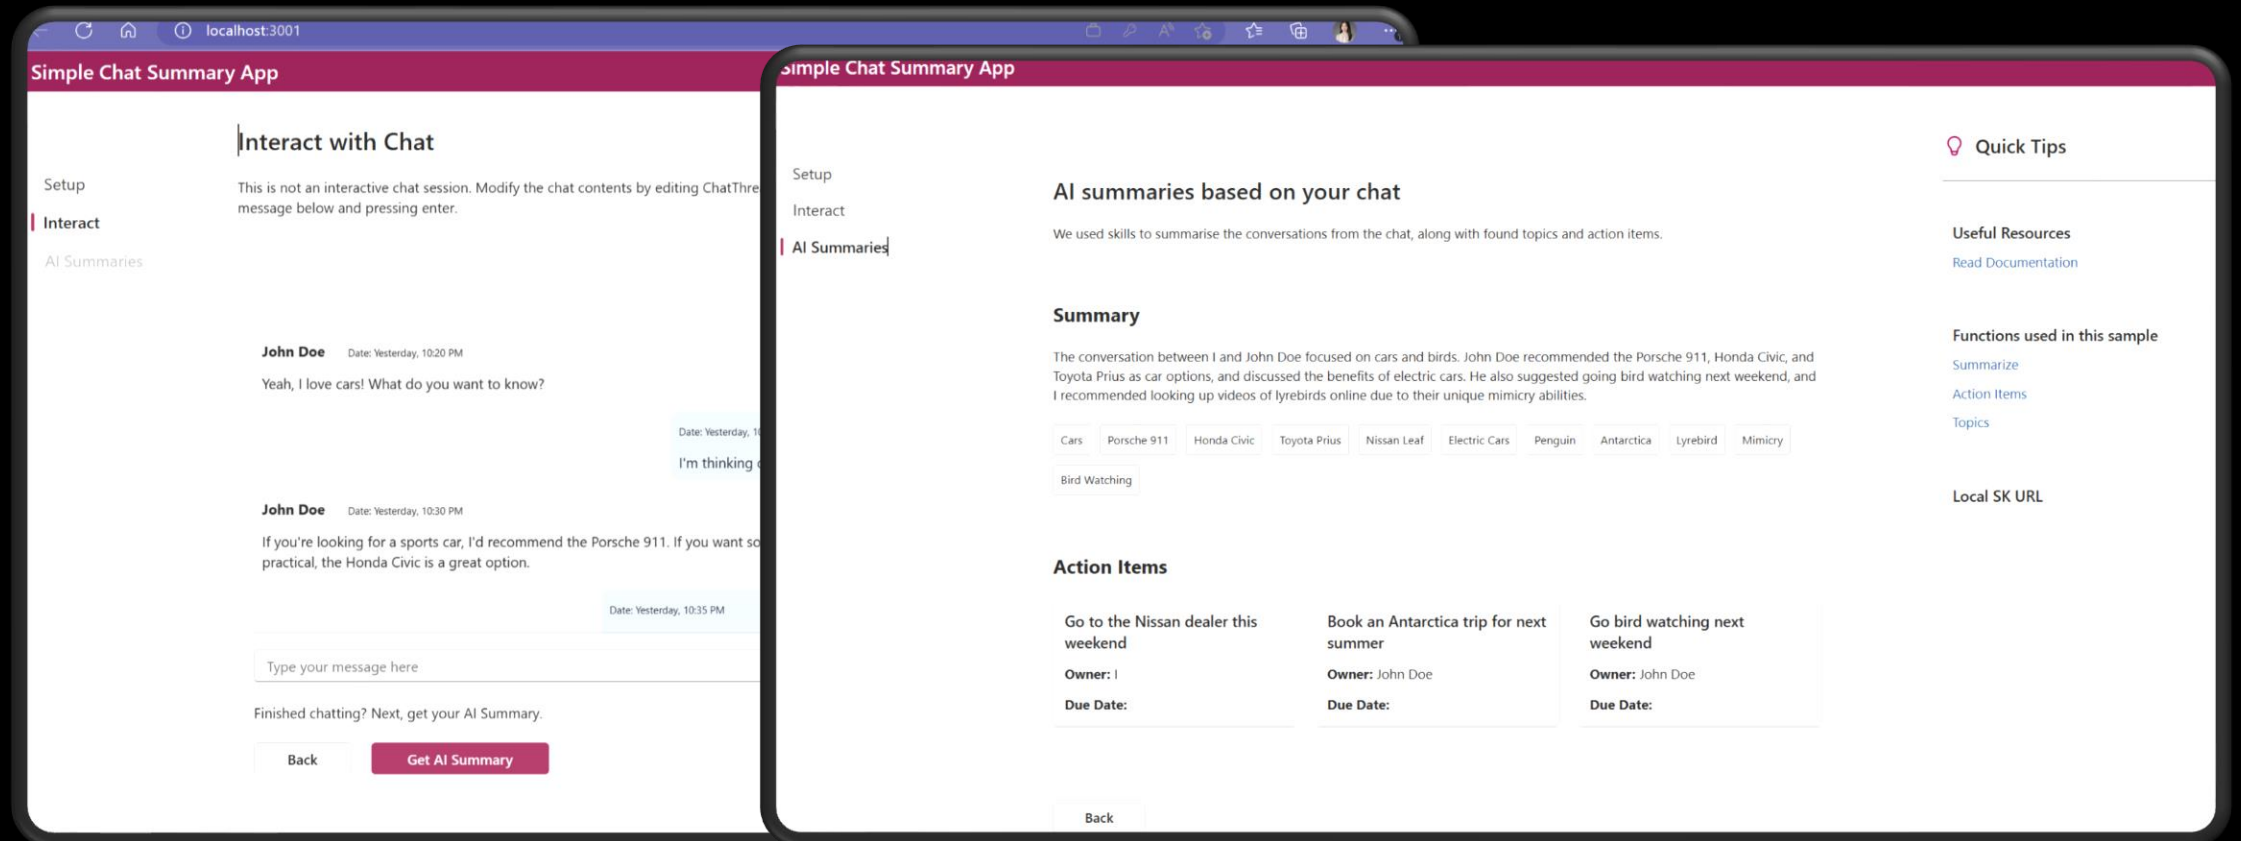
- GitHub Repo Bot** 9:42:54 PM: Hi! I'm your GitHub Repo bot. Here's the repo you are interested in: <https://github.com/microsoft/semantic-kernel>. How can I help you to learn more about this repo?
- User 9:43:26 PM: How Can I start with Github repository Q&A?
- GitHub Repo Bot** 9:43:26 PM: To start with GitHub repository Q&A, you can explore the documentation provided in the repository, such as the `index.md` file in the `semantic-kernel` repository. You can also search for existing questions and answers related to the repository on GitHub, or ask your own questions in the repository's Issues tab.
- User 9:44:12 PM: How can I get started with Semantic Kernel
- GitHub Repo Bot** 9:44:12 PM: To get started with Semantic Kernel, you can read the documentation available in the `dotnet/DocFX/api/index.md` file in the `Microsoft/semantic-kernel` GitHub repository. This file contains

**Input Field:** Type a question you have for the repo

**Right Sidebar:**

- Quick Tips**
- Useful Resources**
  - [Read Documentation](#)
- Functions used in this sample**
  - [Q&A](#)
  - [Pull Web Content](#)
- Learn more about**
  - [Memories](#)
  - [Embeddings](#)
- Local SK URL**

# Simple Chat Summary App



# Workaround(Optional)

samples > apps > github-qna-webapp-react > src > hooks > TS useSemanticKernel.ts > useSemanticKernel

```
1 // Copyright (c) Microsoft. All rights reserved.
2
3 import React from 'react';
4 import { SemanticKernel } from './SemanticKernel';
5
6 export const useSemanticKernel = (uri: string) => {
7     const [semanticKernel] = React.useState(new SemanticKernel('http://localhost:7071'));
8     return semanticKernel;
9 };
10
```

↑  
url



# Q & A

- [Semantic Kernel - Home \(sharepoint.com\)](https://sharepoint.com)
- [Azure-OpenAI-App-Innovation-Workshop/Workshop Content EN at main · xuhaoruins/Azure-OpenAI-App-Innovation-Workshop \(github.com\)](https://github.com/xuhaoruins/Azure-OpenAI-App-Innovation-Workshop)